

Informed Search: A* Algorithm

Algorithmics, 186.814, VU 6.0

Günther Raidl

Algorithms and Complexity Group
Institute of Computer Graphics and Algorithms
TU Wien

WS 2022/23, October 3, 2022




Literature for this Part

- S. Russel, P. Norvig: Artificial Intelligence – A Modern Approach, Algorithm Design, 3rd edition, Prentice Hall, 2009
- J. Pearl: Heuristics: Intelligent Search Strategies for Computer Problem Solving, Addison-Wesley, 1984

Path Planning – Ad Shortest Path Search

- Find a quickest/shortest route $s \rightsquigarrow t$ on a street network




- Which shortest path algorithm to use?
- Only nonnegative weights \rightarrow Dijkstra's Algorithm?
 - Simple implementation: $O(|V|^2)$
 - \rightarrow Typically too slow for large street networks!
- How can we improve the situation?
 - Exploit properties of graph: **sparse, approximately Euclidean**
 - **Apply heuristic information to speed up calculation!**

Dijkstra's Algorithm: Uninformed Search

Dijkstra's algorithm always examines the closest not-yet-examined node.

Example: Simple grid graph



Here, Dijkstra's algorithm corresponds to Breadth-First-Search (BFS).

Informed Search:


- Use an estimation/heuristic function to evaluate each already reached node
- Always expand a most promising node

Greedy Best First Search

Heuristic $h(x)$: “direct” (unobstructed) distance to goal t ,

- e.g. Euclidean distance $d_2(x, t)$ or Manhattan distance $d_1(x, t)$

Always expand a reached node with minimum distance to goal




Here, the least number of nodes is expanded.


BUT: In general, does not guarantee optimal solution!

Greedy Best First Search May Fail

Greedy Best First Search:



Dijkstra's algorithm:



What can we do to avoid such stupid mistakes?

Can we combine the advantages to get a fast exact algorithm?

→ A* algorithm

A* Algorithm – Basic Idea



(Hart, Nilsson, Raphael, 1968)

Priority function: $f(x) = g(x) + h(x)$, $\forall x \in V$, where

$g(x)$: best cost so far to reach x

$h(x)$: estimated cost from x to goal t (“cost-to-go”),
e.g. $d_1(x, t)$ or $d_2(x, t)$

I.e., $f(x)$ estimates the total cost from s to t via node x .




Implementation of A^*

Algorithm 1: A^* Algorithm

```
1 Input: digraph  $G = (V, A)$  with  $w_{ij} > 0 \ \forall (i, j) \in A$ ,  
   source node  $s$ , target node  $t$ ;  
2  $g(s) = 0$ ;  $g(x) = \infty \ \forall x \in V \setminus \{s\}$ ;  
3 priority queue  $Q \leftarrow \{(s, f(s) = h(s))\}$ ;  
4 while  $Q \neq \emptyset$  do  
5    $x \leftarrow Q.\text{getMin}()$  // remove node with min.  $f(x)$ ;  
6   if  $x = t$  then  
7     return  $f(t)$ , path given by predecessor list  $\text{pred}(t)$ ;  
8   for all  $v$  adjacent to  $x$  do  
9      $g' \leftarrow g(x) + w_{xv}$ ;  
10    if  $g' < g(v)$  then  
11       $g(v) \leftarrow g'$ ;  $\text{pred}(v) \leftarrow x$  // new/better path to  $v$ ;  
12       $Q.\text{put}(v, g' + h(v))$   
13 return no path to target  $t$ ;
```

Exemplary A* Run



A* will first expand the nodes s , a , b , and c ,
by which it has found the path (s, a, b, c, t) to the target.
It continues by expanding d ,
finding the shortest path (s, d, t) .


Important:

A* continues until the target node t is selected for expansion.

A*: Re-Expanding Nodes

A* may also find shorter paths to nodes that have already been expanded.

→ Such nodes get re-expanded.



Now, s , a , b , and c are expanded, t reached; $f(c) = 10 + 1 = 11$.
Then d is expanded, yielding a shorter path to c ;
 $f(c) = 9 + 1 = 10$.

A* Terminates

Theorem

A will always terminate in limited time.*

Proof:

- The inner loop performs $\deg(x) = O(|V|)$ iterations, i.e. terminates.
- There are finitely many acyclic paths from s to other nodes in G .
- A* only ever considers acyclic paths.
- On each major iteration a new acyclic path is considered because:
 - When a node is expanded the first time, a new path is considered.
 - When a node is re-expanded, a shorter and thus yet unconsidered route is considered.
- Thus, the most work A* could do is to look at every acyclic path.

A* Always Finds a Path if one Exists

Theorem


If a path from s to t exists, A will always find one.*

Proof:

- Similarly as DFS and BFS:
- After expanding a node x , A* has found paths from s to all nodes directly reachable via x .
- Only case when A* stops without having found a path to t :
 - When no further reachable node exists that has not yet been expanded.
 - This, however, implies that no path to t exists.

Is A^* Guaranteed to Find an Optimal Path?

No! At least not as introduced so far:



After expanding s , t will be selected for expansion and the algorithm stops with the path (s, t) .

Under which conditions is A^* guaranteed to find an optimal path?

Admissible Heuristics $h(x)$

Let $h^*(x)$ be the true minimal cost from x to t .

Definition (Admissible heuristic)

A heuristic $h(x)$ is **admissible** if it represents a **lower bound**, i.e.

$$h(x) \leq h^*(x) \quad \forall x \in V.$$

- An admissible heuristic is “optimistic”.
- In the grid-graph examples, $d_1(x, t)$ and $d_2(x, t)$ are admissible heuristics.
- $h(x) = 0$ also is an admissible heuristic.

A^* with an Admissible Heuristic Yields an Optimal Path

Theorem

If $h(x)$ is admissible, A^ is guaranteed to find a least-cost path.*

Proof:

- Suppose it finds a suboptimal path with cost $f' > f^* = h^*(s)$.
- There must exist a node x which is
 - unexpanded
 - and whose path from s (as provided by pred) is the beginning of a true shortest path to t .
- $f(x) \geq f'$ (else A^* would not have terminated), and

$$\begin{aligned}f(x) &= g(x) + h(x) \\&= g^*(x) + h(x) \\&\leq g^*(x) + h^*(x) \\&= f^*\end{aligned}$$

- So, $f^* \geq f(x) \geq f'$, contradicting the assumption.

Monotonic Heuristics

Definition (Monotonic heuristic)

A heuristic $h(x)$ is **monotonic (consistent)** if

$$h(x) \leq w_{xy} + h(y) \quad \forall (x, y) \in A \quad \text{and} \quad h(t) = 0$$

- Monotonic heuristics ensure that for any path X from s to x

$$f(x) = g(x) + h(x) \leq g(x) + w_{xy} + h(y) = g(y) + h(y).$$

- I.e., it is impossible to decrease $f(x)$ by extending a path to include a neighboring node.
- **No re-expansions of nodes.**
- Comparable with the situation of non-negative edge weights in Dijkstra's algorithm.
- **Monotonicity implies admissibility of the heuristic!**

Monotonic Heuristics (contd.)

- Are $d_1(x, t)$ and $d_2(x, t)$ monotonic in our grid-graph examples?
 - Yes.
- For street networks finding a reasonable monotonic/admissible heuristic
 - is typically easier for finding **shortest** routes (e.g. Euclidean distance),
 - but not so for **fastest** routes.

Further Properties of A^*

- A^* is optimally efficient w.r.t. the used heuristic $h(x)$, meaning that no complete algorithm employing the same heuristic will expand fewer nodes (proof omitted).
- Dijkstra's algorithm can be viewed as the special case of A^* with $h(x) = 0$.

Runtime:

- in general $O(d^l)$, where l is the number of nodes on the shortest path, d a constant, but actual runtime depends strongly on $h(x)$
- better heuristic \rightarrow better A^* performance:
if $h_1(x) \leq h_2(x) \forall x \in V \rightarrow h_2$ dominates (or is equal to) h_1
 $\rightarrow A^*$ is guaranteed to expand no more nodes with h_2 than with h_1
- with an (almost) perfect heuristic $h(x) \approx h^*(x)$ runtime can be $\Theta(l)$

Example: 8-Puzzle

Example State

1		5
2	6	3
7	4	8

Goal State

1	2	3
4	5	6
7	8	

What heuristic to use?

- $h_0(x) = 0$
- $h_1(x)$ = number of tiles in wrong positions
- $h_2(x)$ = sum of Manhattan distances of each tile to its target location


# moves, i.e., f^*	$A_{h_1}^*$ -nodes	$A_{h_2}^*$ -nodes	$A_{h_1}^*$ -bf	$A_{h_2}^*$ -bf
6	20	8	1,33	1.24
14	539	113	1,44	1,23
24	39.135	1.641	1,48	1,26

bf: branching factor, i.e., avg. number of exp. successors (Russel, Norvig; 2009)

Further Improvements/Variants

For complex problems, A^* sometimes still too slow and/or needs too much memory. Further improvement possibilities:

- **Bidirectional search:** search from s to t and t to s at the same time



- **Weighted A^* :** “Inflate” heuristic, i.e. put more emphasis on $h(x)$
 - not admissible anymore, i.e. optimality not guaranteed
 - however, it might be possible to obtain a bound on sub-optimality (approximation guarantee)
- **Preprocessing of data and store information for improved heuristic:**
 - Extrem case: Calculate all-pairs shortest paths, store pairwise minimum distances, i.e. $h^*(x)$

Further Improvements/Variants (contd.)

- **Hierarchical approach:** store graph and compute shortest paths on a hierarchy of abstraction levels
 - E.g. street network:
 - First, plan rough route only considering major motor ways
 - Then refine, considering main roads connecting cities/villages
 - Finally, consider all smaller streets
- **Iterated Deepening A* (IDA):** primarily to limit memory usage
 - Consider integer costs; for costs $c = 0, 1, \dots$ do
 - Perform depth-first search not expanding any node x with $f(x) > c$.
 - If t expanded, return path to it.
 - Guaranteed to find optimal solution but in general much slower than A*.